

SIGEVO Summer School (S3)

Modelling Optimization Problems for Evolutionary Algorithms and Other Metaheuristics

Project statement

Carlos M. Fonseca
Department of Informatics Engineering
University of Coimbra
Portugal

6 July 2021

1 Introduction

This project consists in modelling given combinatorial optimization problems as local search problems and implementing the corresponding models while bearing in mind the principles and ideas presented in the lecture. Problem implementations shall adhere to an Application Programming Interface (API) so that a number of simple algorithms, or *solvers* (provided together with the API specification), can run directly on them in a problem-independent way, and experimental results can be collected.

Four different problems are proposed. Each problem should be modelled and implemented either in C (the preferred language) or in Python3 (for two of the four problems only) by a group of 3 to 5 students (typically 4). It is essential that most, but not necessarily all, members of each group are comfortable programming (all in the same language), so that coding work can be effectively shared during the time available. Group members who do not program should still be able to contribute actively to other aspects of the project work, including conceptual modelling, experimentation, and analysis and discussion of results.

It is expected that each group gives a presentation on the last day of the School. Presentations should cover a description of the problem considered, modelling decisions, such as solution representation and neighbourhood structure, implementation decisions, including search space sampling, search operators, and objective function evaluation strategies, and (preliminary) experimental results.

1.1 Software requirements

The required development environment depends on the language of choice:

C A modern C compiler such as a recent version of `gcc` or `clang` is required. The ability to compile against the GNU Scientific Library (v2.5 or later) is also required, although this requirement can be worked around in case it proves difficult to meet.

In addition, development tools such as `make` (build automation), a debugger, and `valgrind` (dynamic analysis) will be very useful, but are not strictly required.

Python3 A Python3 installation containing the numerical package Numpy is required.

Students should ensure that they have a working programming environment before the School starts.

1.2 Programming skills

The C API uses Abstract Data Types (ADTs) implemented as pointers to structures. Students should be familiar with structures, arrays, and pointers in C.

Python3 implementations will consist of concrete classes that derive from a provided set of abstract classes, and Numpy arrays are used throughout. Students should be familiar with object-oriented programming in Python and with Numpy. Numpy array arithmetic, indexing and universal functions (“ufuncs”) are essential for the problems modelled in Python to achieve reasonable performance.

2 Problem descriptions

2.1 3D printing service

Programming language: Python or C

An additive manufacturing startup provides 3D printing services both to private and business customers. A single 3D printer is used to produce custom metal parts that must be printed without interruption. Once a part is printed, the machine becomes available after a fixed change-over time. Each part must be ready for delivery by a given deadline, which is set at the time the order is placed. Missing the deadline leads to a penalty that depends on the part, and is proportional to how late the part is completed. The time needed to print a given part can be determined in advance.

The problem consists in determining the order in which a set of parts should be manufactured on the single metal 3D printer so as to minimize the total penalty incurred by the service provider.

2.2 Urban waste collection

Programming language: Python or C

Waste collection vehicles are used to regularly empty the many public waste containers located on the side of roads (away from intersections) across Neat City. A given city quarter is served by a single vehicle that is large enough to collect the waste from all containers in a single tour. A collection tour consists in the vehicle leaving the depot, visiting and emptying all containers in the quarter, disposing of the collected waste at the city's waste treatment plant, and returning to the depot.

On one-way and most two-way roads, the vehicle may empty the containers it passes by regardless of whether or not it is driving on the side of the road where the container is located. On multi-lane roads, however, the vehicle and the container must be on the same side of the road for the collection to be possible. In any case, no U-turns are allowed at waste containers.

Suppose that the cost of a waste collection tour is proportional to the total distance travelled by the vehicle until it returns to the depot. The problem consists in determining the shortest tour subject to the above one-way, roadside, and no U-turn constraints.

2.3 Campus network

Programming language: C

New network infrastructure is to be installed at a University campus. Each building on campus will be connected directly to the computing centre by a dedicated high-speed cable. Trenches will be dug in order to allow cables to be laid from one building to another. Once dug, a trench may be used to lay more than one cable, such as those connecting other buildings farther away from the computing

centre. Note that trenches may only be dug from one building to another, and that different trenches may not cross each other.

The cost of digging a trench is proportional to its length, as is the cost of a cable. The problem consists in determining what trenches to dig so as to minimize the total (trench and cable) set up cost of the new network.

2.4 Community detection

Programming language: C

A fully-connected undirected graph, where vertices represent users and weighted edges represent the intensity of some attribute of their interaction that may be positive or negative, was obtained from social network data. Users connected by edges with positive weight show affinity to each other, whereas negative edge weights indicate lack of affinity. Groups of users connected mostly by positively weighted edges suggest the existence of a community involving those users.

The problem consists in finding the groups of users that maximize the total internal edge weight of all groups.

3 Problem modelling

Modelling each of the above problems as a local search problem begins with attempting to answer the following questions:

Problem instance What (known) data is required to fully characterize a particular *instance* of the problem? This must be available in advance, and is not changed by the solver in any way.

Solution What (unknown) data is required to fully characterize a (feasible) candidate *solution* to a given problem instance? This is the data needed to implement the solution in practice, and will be determined by the solver during the optimization run.

Objective function How can the performance of a given candidate solution be measured? This depends only on the problem instance and the actual solution itself, and never on how the solution was actually found. Is the corresponding value to be minimized or maximized?

Neighbourhood structure What makes two given solutions *similar* to each other? This usually means that parts of the two solutions are somehow identical, but it should also happen that they exhibit *similar* performance *in most cases*. A candidate solution that performs at least as well as all solutions

similar to it (its *neighbours*) is called a *local optimum* of the corresponding problem instance.

The choice of the neighbourhood structure is particularly important for the success of local search algorithms. By ensuring that similar solutions tend to exhibit similar performance, one seeks to induce fewer local optima and large basins of attraction to those optima, although this can seldom be guaranteed. Furthermore, any two feasible solutions should be connected by a sequence of consecutive neighbours, so that the unknown global optimum can, at least in principle, be reached from any initial solution.

3.1 Computational model

Once suitable answers to the above questions are obtained, a more refined set of questions relative to the computer implementation of the model can be considered.

Problem instance representation How should the problem instance data be stored in a data structure, so that the objective function can be easily computed?

Solution representation How should candidate solutions be represented, i.e., stored as a data structure, so that:

1. Their performance can be evaluated efficiently through the objective function?
2. They can be easily modified to obtain neighbouring solutions?

Solution evaluation How can the objective function be computed given the instance data and the solution representation?

Move representation How can *moves*, i.e., changes that, when applied to a solution, lead to a neighbouring solution, be represented?

Incremental solution evaluation When an evaluated solution is modified by applying one or more moves to it, can the resulting solution be evaluated faster than would otherwise be the case? How?

Move evaluation How much would applying a given move to a solution change its performance? Can this effect be computed more efficiently without modifying the original solution than by evaluating it, applying the move and evaluating the result? How?

3.2 Search operators

Having established how solutions and moves are represented and evaluated, the problem model needs to provide the elementary problem-dependent *operations* used by local search solvers to explore the search space. The following few operations are commonly used by evolutionary algorithms and other local-search metaheuristics:

Random solution generation is typically used to obtain (a diverse set of) initial solutions to start the search. In the absence of instance-specific inside knowledge about the location of the optimum, this generation is ideally performed *uniformly* at random. Since the search space is typically large, the probability of generating the same solution twice when sampling with replacement should be negligible.

Random move generation (with replacement) arises in perturbation operators, such as mutation in EAs, whenever neighbourhood exploration is limited to a small number of random neighbours. It is usually simpler and more efficient to implement than sampling without replacement.

Random move generation (without replacement) is required by solvers that potentially perform a complete exploration of the search neighbourhood in order to identify local optima, such as Iterated Local Search.

More elaborate operations, including recombination, differential evolution, and particle-swarm operators, can also, at least in principle, be implemented in terms of suitable sampling operations based on the specified neighbourhood structure. Geometric recombination operators, for example, can be constructed based on the random generation of moves from one parent solution *towards* the other, in the sense of edit distance between solutions. However, such operators are beyond the scope of this project.

4 Application Programming Interfaces

The symbols marking the items of these interfaces have the following meaning:

- Basic elements required to support simple random search.
- ▷ Additional elements required to support a simple mutation-only evolutionary algorithm.
- ◇ Further elements required to support an iterated local search algorithm.

4.1 C API

This is (a subset of) the `nasf4nio` API with a single modification related to problem instantiation. In your project, you should use the provided `s3problem.h` header file instead of `problem.h`.

4.1.1 Data structures

- `struct problem {...}`
- `struct solution {struct problem *p;...}`
- ▷ `struct move {struct problem *p;...}`

4.1.2 Problem instantiation and inspection

- `struct problem *newProblem(char *filename) {...}`
Allocate a problem structure and initialize it with the problem-specific data contained in the specified input file.
- `int getNumObjectives(const struct problem *p) {return 1;}`

4.1.3 Memory management

- `void freeProblem(struct problem *p) {...}`
Free all memory used by a problem structure.
- `struct solution *allocSolution(struct problem *p) {...}`
Allocate memory for a solution.
- `void freeSolution(struct solution *s) {...}`
Free all memory used by a solution structure.
- ▷ `struct move *allocMove(struct problem *p) {...}`
Allocate memory for a move.
- ▷ `void freeMove(struct move *v) {...}`
Free all memory used by a move structure.

4.1.4 Reporting

- `void printProblem(struct problem *p) {...}`
Print a user-formatted representation of a problem instance.

- `void printSolution(struct solution *s) {...}`
Print a user-formatted representation of a solution.
- ▷ `void printMove(struct move *v) {...}`
Print a user-formatted representation of a move.

4.1.5 Operations on solutions

- `struct solution *randomSolution(struct solution *s) {...}`
Uniform random sampling of the solution space.
The input argument must be a pointer to a solution previously allocated with `allocSolution()`, which is modified in place.
If `randomMoveWOR()` is implemented, this function must also initialize the corresponding state by performing the equivalent to `resetRandomMoveWOR()`.
The function returns its first input argument.
- `struct solution *copySolution(struct solution *dest, const struct solution *src) {...}`
Copy the contents of the second argument to the first argument, which must have been previously allocated with `allocSolution()`. The function returns its first input argument.
- `double *getObjectiveVector(double *objv, struct solution *s) {...}`
Single or multiple objective full and/or incremental solution evaluation.
Once a solution is evaluated, results may be cached in the solution itself so that a subsequent call to this function simply returns the precomputed value and/or the solution can be re-evaluated more efficiently after it is modified by one or more calls to `applyMove()`. Therefore, the formal argument is not `const`. Solution (re-)evaluation must occur before this function returns, but the time at which actual evaluation occurs is otherwise left unspecified.
The function returns its first input argument.
- ▷ `struct solution *applyMove(struct solution *s, const struct move *v) {...}`
Modify a solution by applying a move to it. It is assumed that the move was generated for, and possibly evaluated with respect to, that particular solution. The result of applying a move to a solution other than that for which it was generated/evaluated (or a pristine copy of it), including applying the same move to a solution more than once, is undefined.
If `randomMoveWOR()` is implemented, this function must also initialize the corresponding state by performing the equivalent to `resetRandomMoveWOR()`.
The function returns its first input argument.

- ◇ `int getNeighbourhoodSize(struct solution *s) {...}`
Return the number of neighbours of a given solution.
- ◇ `struct solution *resetRandomMoveWOR(struct solution *s) {...}`
Reset the uniform random sampling without replacement of the neighbourhood of a given solution, so that any move can be generated by the next call to `randomMoveWOR()`. The function returns its input argument.

4.1.6 Operations on moves

- ▷ `struct move *randomMove(struct move *v, const struct solution *s) {...}`
Uniform random sampling of the neighbourhood of a given solution, with replacement. The first input argument must be a pointer to a move previously allocated with `allocMove()`, which is modified in place.
- ▷ `struct move *copyMove(struct move *dest, const struct move *src) {...}`
Copy the contents of the second argument to the first argument, which must have been previously allocated with `allocMove()`. The function returns its first input argument.
- ◇ `double *getObjectiveIncrement(double *obji, struct move *v, struct solution *s) {...}`
Single or multiple objective move evaluation with respect to the solution for which it was generated, before it is actually applied to that solution (if it ever is). The result of evaluating a move with respect to a solution other than that for which it was generated (or to a pristine copy of it) is undefined. Once a move is evaluated, results may be cached in the move itself, so that they can be used by `applyMove()` to update the evaluation state of the solution more efficiently. In addition, results may also be cached in the solution in order to speed up evaluation of future moves. Consequently, neither formal argument is `const`. The function returns its first input argument.
- ◇ `struct move *randomMoveWOR(struct move *v, struct solution *s) {...}`
Uniform random sampling of the neighbourhood of a given solution, without replacement. The first input argument must be a pointer to a move previously allocated with `allocMove()`, which is modified in place. The function returns this pointer if a new move is generated or `NULL` if there are no moves left.

4.2 Python3 API

4.2.1 Problem class

A problem is implemented as a single class that defines two inner classes and a number of “standard” methods. The high level structure of a problem class definition is as follows:

```
import numpy as np
import sample

class ProblemX:
    class Solution(sample.Element):
        ...
    class SolutionSample(sample.Sample):
        ...
    class Move(sample.Element):
        ...
    class MoveSample(sample.Sample):
        ...
    def __init__(self, ...):
        ...
    def randomSolution(self, n=None):
        if n is None:
            # return a single solution
            return self.Solution(self, ...)
        else:
            ...
            # return a sample of n solutions
            return self.SolutionSample(self, ...)
```

Classes `Solution` and `Move` are used to store single solutions and single moves, respectively. Classes `SolutionSample` and `MoveSample` are containers for sequences of solutions and moves, respectively, and should emulate part of the behaviour of 1-dimensional Numpy arrays. In particular, indexing a `SolutionSample` object with an integer should return a `Solution` object. `Solution` objects can be transformed into `SolutionSample` objects via a `.toSample()` method, and analogously for `Move` and `MoveSample` objects.

The `__init__` method implements problem-specific instantiation from user-supplied instance data, which is typically stored as attributes of the newly instantiated problem instance.

The `randomSolution` method implements random solution generation. Depending on the value of `n`, the generated solutions are returned as a newly instantiated `Solution` or `SolutionSample` object, as appropriate.

4.2.2 Solution classes

The `Solution` class should have the following structure:

```
class Solution(sample.Element):
    # argument 'of' is the outer problem class object
    def __init__(self, of, data, ...):
        self.of = of
        self.sample = of.SolutionSample
        ...
    # make a copy of the current solution
    def copy(self):
        return self.__class__(self.of, ...)
    # toSample() should return a "view" of the current
    # solution data, not a copy
    def toSample(self):
        return self.sample(self.of, ...)
    # apply single move 'other' to current solution, in place
    def __iadd__(self, other):
        ...
    # apply move(s) to current solution
    def __add__(self, other):
        ... # return a Solution if 'other' is a Move object, or a
            # SolutionSample object if it is a MoveSample object
    # random move generation
    def randomMove(self, n=None):
        if n is None:
            ... # return a Move object
        else:
            ... # return a MoveSample object
    # objective function evaluation
    def objvalue(self):
        ... # return a scalar
```

The `SolutionSample` class should implement essentially the same methods as the `Solution` class, except `toSample`. Additional methods and any differences with respect to the corresponding `Solution` methods are detailed next:

```

class SolutionSample(sample.Sample):
    ...
    # apply move to current solution
    def __add__(self, other):
        ... # return a SolutionSample
    def randomMove(self, n=None):
        ... # return a MoveSample object
    # emulate 1-d array indexing/slicing
    def __getitem__(self, key):
        ...
    def __setitem__(self, key, value):
        ...
    # emulate the Numpy array repeat() method
    def repeat(self, n):
        return self.__class__(self.of, ...)
    # apply moves to current solution
    def __add__(self, other):
        ... # return a SolutionSample object
    # random move generation
    def randomMove(self, n=None):
        ... # return a MoveSample object
    # objective function evaluation
    def objvalue(self):
        ... # return a 1-dimensional Numpy array

```

4.2.3 Move classes

The Move and MoveSample classes can usually be implemented as follows, where data should contain the representation of the move(s).

```

class Move(sample.Element):
    def __init__(self, of, data):
        self.of = of
        self.sample = of.MoveSample
        self.data = np.asarray(data)

class MoveSample(sample.Sample):
    def __init__(self, of, data):
        self.of = of
        self.element = of.Move
        self.data = np.asarray(data)

```

In this case, the default methods inherited from the parent classes do not need to be redefined.

4.2.4 Additional remarks

In practice, there is usually no need implement the same method separately for both `Solution` and `SolutionSample`. In fact, provided that a given method is implemented for the latter, it can also be easily implemented for the former by promoting the solution to a sample using `.toSample()`, calling the desired method, and indexing the result to obtain a solution again. Alternatively, if the method is initially implemented for `Solution` only, a simple implementation for `SolutionSample` can be produced by iterating over `self` and successively calling the existing `Solution` method.

One last feature concerns operations involving solution samples and move samples of different sizes. Calling `randomMove` with a value of `n` greater than 1 will generate a move sample with n times the elements of the original solution sample. In this case, the first n moves correspond to the first solution, the next n moves correspond to the second solution, and so on. The `__add__` methods should operate in accordance to this.

The supplied implementation of the n -queens problem illustrates many of these aspects, and can be used as a reference.